

A Reproducible Tutorial on Reproducibility in Database Systems Research

VLDB 2024

Denis Hirn

{denis.hirn, tim.fischer}@uni-tuebingen.de

University of Tübingen

Tim Fischer

Gökhan Kul

gkul@umassd.edu

University of Massachusetts

Overview



Part 1

Reproducibility Definition and Goals

Part 2

Containerization Technologies

Part 3

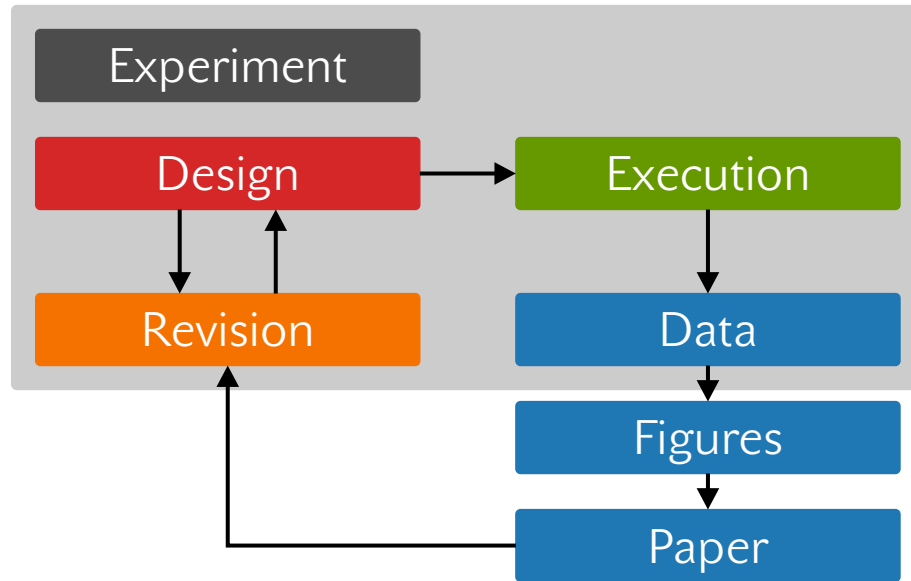
Automation

Part 4

Documentation



Reproducibility Workflow in Database Systems Research



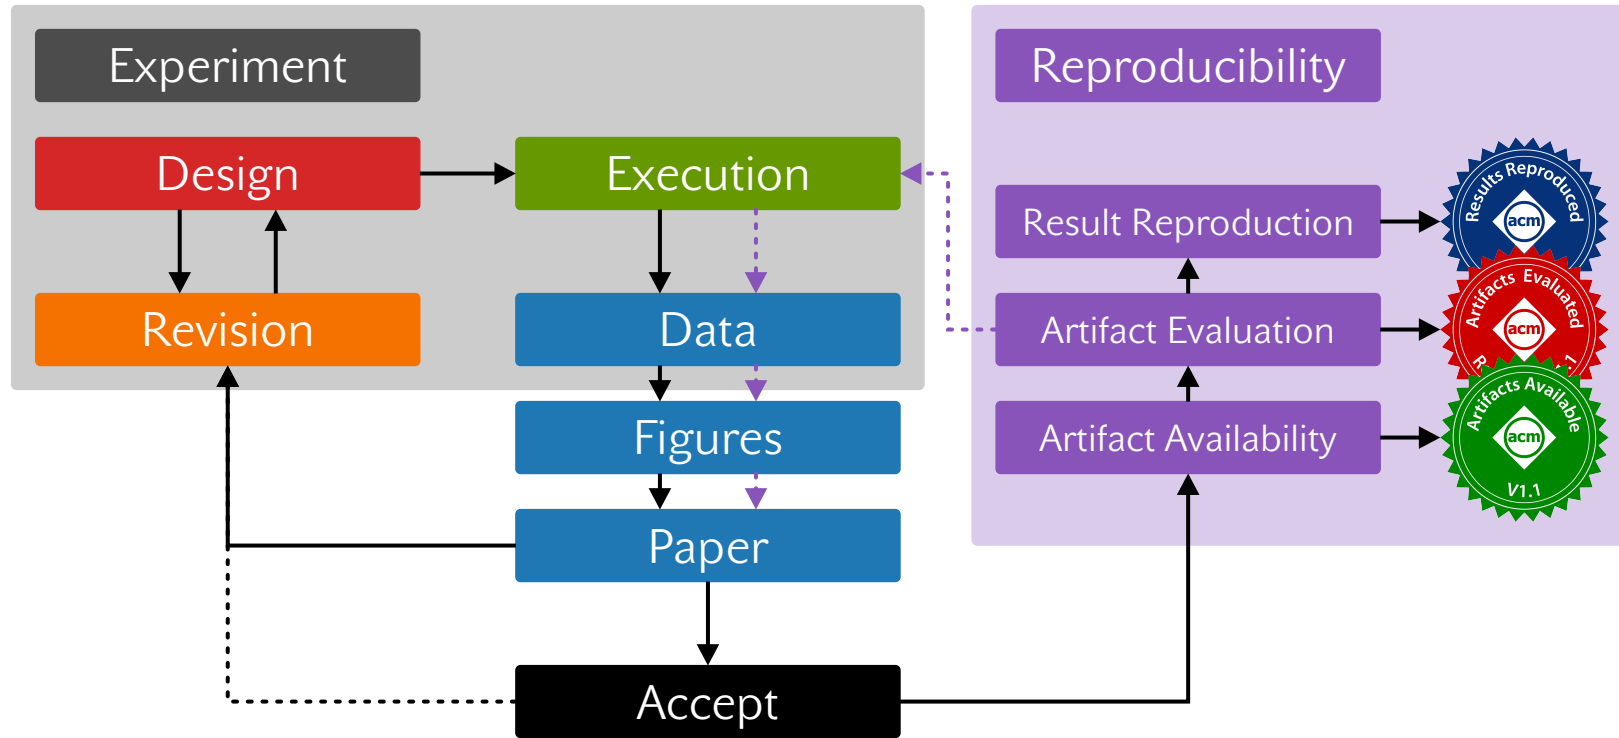
SIGMOD reproducibility effort:

<https://reproducibility.sigmod.org>

VLDB reproducibility effort:

<https://vldb.org/pvldb/reproducibility>

Reproducibility Workflow in Database Systems Research



SIGMOD reproducibility effort:

<https://reproducibility.sigmod.org>

VLDB reproducibility effort:

<https://vldb.org/pvldb/reproducibility>



Goals of Reproducibility

Availability

- ▶ The artifacts (software, data, and documentation) should be available.

Reproducibility

- ▶ Central results and claims should be supported.
- ▶ Experimental results can be independently verified.
- ▶ Exact replication is not always possible, but the overall results should match conclusions.

In short: More impact, more trust, more credibility.



Reproducibility in Practice

Reproducibility is...


- ▶ the ability to re-run an experiment and obtain the same results,
 - ▶ a cornerstone of the scientific method,
 - ▶ a key factor in the *credibility* of research results.
-

However...

- ▶ it can be difficult to achieve in practice,
- ▶ especially in database systems research,
- ▶ and when done as an afterthought.

 The key to reproducible experiments is to design them with reproducibility in mind.

Containerization Technologies



docker



Without Containers

- ▶ Ship packages: `zip`, `tar.gz`, etc.
- ▶ Manual installation of software and manual configuration.
- ▶ Dependency hell.
- ▶ “Works on my machine. -_(ツ)_/ -”
- ▶ Experiments are difficult to reproduce, because
 1. the environment (configuration) may be different,
 2. package versions may be different,
 3. the software may not be available anymore
 4. :



In short, it is a **mess**!




With Containers

- ▶ ~~Ship packages: zip, tar.gz, etc.~~
 - ▶ ~~Manual installation of software and manual configuration.~~
 - ▶ ~~Dependency hell.~~
 - ▶ ~~"Works on my machine. -_(ツ)_/-"~~
 - ▶ ~~Experiments are difficult to reproduce, because~~
-
- ▶ Containers contain **everything**. No more packages!
 - ▶ Isolation.
 - ▶ Works everywhere.
 - ▶ Remains functional in the future.
 - ▶ Experiments are easy to reproduce. For future you, and others!



Reproducibility is not a *one-time* effort.

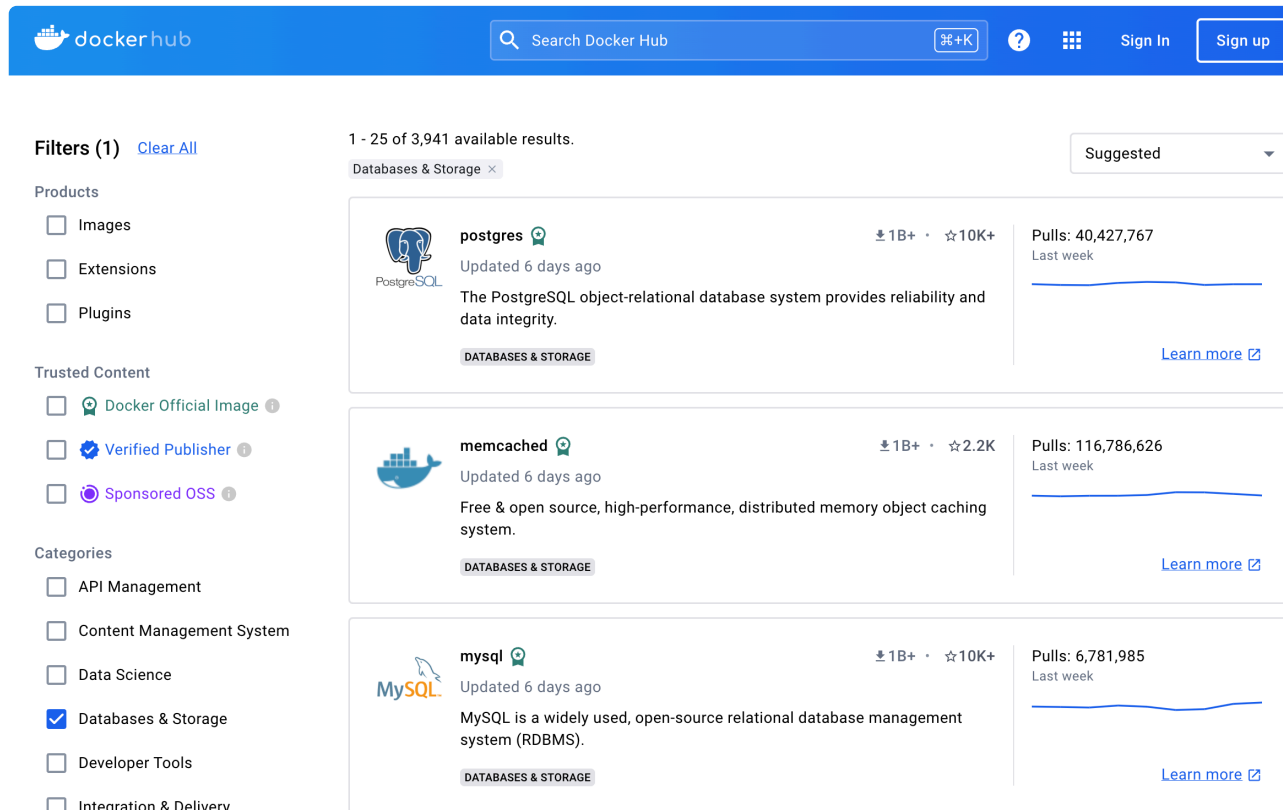
Docker Abstraction Layers



- ▶ The **containers** are isolated from the **host system** and other **containers**.
- ▶ Eliminate the “*works on my machine*” problem.
- ▶ Prevents pollution of the **host system**.



- ▶ Huge repository of pre-built Docker images.
- ▶ Can be used to host your own images (aka. reproducibility artifacts).



The screenshot shows the Docker Hub search results for the 'Databases & Storage' category. The page features a blue header with the Docker Hub logo, a search bar, and navigation links. On the left, there are filters for Products (Images, Extensions, Plugins), Trusted Content (Docker Official Image, Verified Publisher, Sponsored OSS), and Categories (API Management, Content Management System, Data Science, Databases & Storage, Developer Tools, Integration & Delivery). The main content area displays three search results, each with a logo, name, description, and pull statistics.

Image Name	Description	Pulls (Last Week)
postgres	The PostgreSQL object-relational database system provides reliability and data integrity.	40,427,767
memcached	Free & open source, high-performance, distributed memory object caching system.	116,786,626
mysql	MySQL is a widely used, open-source relational database management system (RDBMS).	6,781,985



Dockerfiles and Docker Images



```
1 FROM ubuntu:22.04 AS duckdb
2
3 RUN apt-get update && apt-get install -y \
4     build-essential cmake git
5
6 RUN git clone --depth 1 --branch v1.0.0 \
7     https://github.com/duckdb/duckdb.git
8
9 WORKDIR /duckdb
10 RUN make release -j
11
12 # Add the binary to the PATH
13 ENV PATH="/duckdb/build/release:${PATH}"
14 # Define the entrypoint
15 ENTRYPOINT ["duckdb"]
```

01_simple/Dockerfile.duckdb

```
1 > docker build -t duckdb .
```



Docker Compose



```
1 services:
2   duckdb:
3     image: duckdb
4     build:
5       context: .
6       dockerfile: Dockerfile
7   postgres:
8     image: postgres:16
```

 01_simple/docker-compose.yml

```
1 > docker compose up -d
```



- ▶ Docker Compose is used for multi-container applications.
- ▶ Single command to **start | stop | remove** all services.
- ▶ Simplifies wiring up the containers and volumes.

Docker Volumes



- ▶ **Volumes** are a way to persist data in Docker containers.
- ▶ Can be used to share data between containers.
- ▶ Bind mounts can be used to share data between the host and the container.

```
1 services:
2   duckdb:
3     image: duckdb
4     build:
5       context: .
6       dockerfile: Dockerfile
7     volumes:
8       - type: bind
9         source: ./data
10        target: /data
11   postgres:
12     image: postgres:16
13     volumes:
14       - postgres_volume:/var/lib/postgresql/data
15       - type: bind
16         source: ./data
17        target: /data
```

 01_simple/docker-compose.yml

```
1 |> docker compose up -d
```





Takeaways from Containerization



Containers provide consistent software environments for experiments.



Images can be shared, reused, and versioned.



Containers eliminate manual software installation on the host machine.

Automation



GNU Make



Automate as Much as Possible

- ▶ Automation is key to reproducibility.
- ▶ Experimental execution should be as simple as possible.
- ▶ The user should not have to worry about the details.
- ▶ The user should be able to reproduce the results with a single command.
- ▶ Post-processing should be automated, too.



Experimental Scenario

- ▶ Run a set of experiments (*e.g.*, TPC-H queries)
 1. On DuckDB
 2. On PostgreSQL
- ▶ Goal is to compare the performance of the two systems.



Automation Script

```
1 mkdir -p ./data
2
3 docker compose build
4
5 docker compose up postgres -d --wait
6
7 docker compose run --rm duckdb bash "scripts/generate-tpch-duckdb.sh"
8
9 docker compose run --rm duckdb bash "scripts/run-experiments-duckdb.sh"
10
11 docker compose exec postgres bash "scripts/load-tpch-postgres.sh"
12
13 docker compose exec postgres bash "scripts/run-experiments-postgres.sh"
14
15 docker compose down -v --rmi all
```

 01_simple/run.sh

-
- ▶ Single “entrypoint” for the user.
 - ▶ Include setup, execution, and cleanup.



Auditability

```
1 #? █
2 :
3 ⌚ however long later
4 :
5 Done.
6 > █
```

⊗ No indication of progress, about what is happening, or how long it will take.

```
1 #? █
2 :
3 Running experiment 1...
4 ...
5 3 / 5:
6 :
7 Done.
8 > █
```

✓ Output the progress of the experiments to the console.

- ▶ Auditability is the ability to trace the execution of a script.
- ▶ Scripts should provide feedback to the user.
- ▶ Progress indicators can be helpful.
- ▶ Log files should be timestamped and written to disk.
- ▶ Estimating the time to completion can be helpful.



Dockerized TeX Compilation



```
1 services:
2   # ...
3   latex:
4     image: danteev/texlive
5     volumes:
6       - type: bind
7         source: .
8         target: /project
9     working_dir: /project
```

 02_intermediate/docker-compose.yml

```
1 #!/bin/bash
2
3 mkdir -p /project/data/paper
4 pushd /project/paper > /dev/null
5 latexmk\
6   -outdir=/project/data/paper \
7   -pdf \
8   -interaction=nonstopmode \
9   -shell-escape \
10  -synctex=1 \
11  -f \
12  paper.tex
13 popd > /dev/null
```

 02_intermediate/scripts/build-paper.sh



Automation using a Build System: make

```
1 .PHONY: docker
2 docker: # Build and start Docker services
3     ...
4
5 .PHONY: setup
6 setup: docker # Setup the environment
7     ...
8
9 .PHONY: experiments
10 experiments: setup # Run the experiments
11     ...
12
13 .PHONY: paper
14 paper: experiments # Build the paper
15     ...
16
17 .PHONY: clean
18 clean: # Clean up everything
19     ...
```

03_advanced/Makefile

-
- ▶ Use `make` to automate the entire workflow.
 - ▶ Each target depends on the previous target.
 - ▶ The `clean` target cleans up everything.



Use Docker within Docker

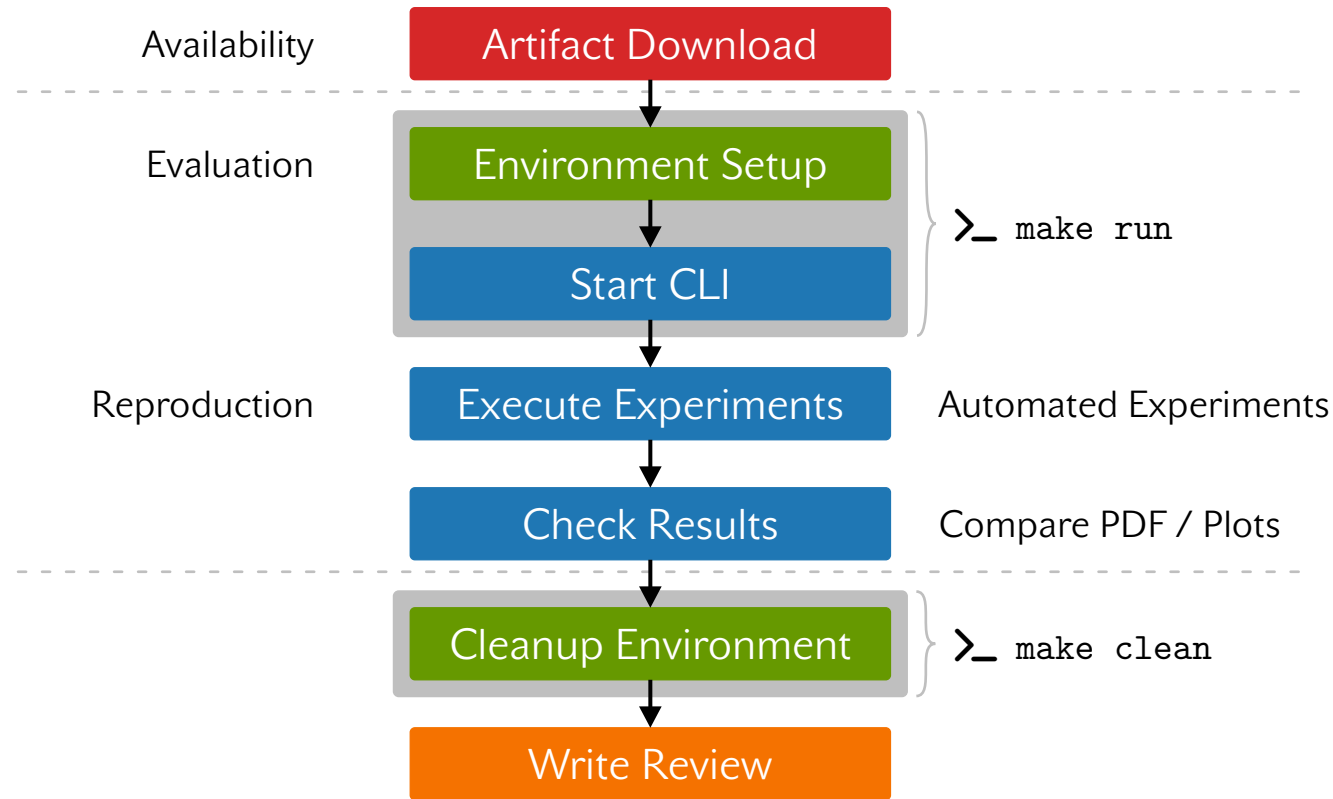
```
1 services:
2   docker:
3     image: runner
4     privileged: true
5     build:
6       context: .
7       dockerfile: Dockerfile
8     volumes:
9       - ./project:/project
10    working_dir: /project
```

04_full/docker-compose.yml

-
- ▶ Docker in Docker is a common pattern for CI/CD pipelines.
 - ▶ Prevents polluting the host system.
 - ▶ Gives scripts full control over the Docker environment, without needing to rely on the host system's environment.
 - ▶ Use the `privileged` flag to run Docker within Docker.



Reproducibility Review Workflow



The review process should be as easy and automated as possible.
It is the authors' responsibility to provide an easily reproducible artifact.



Takeaways from Automation



1 Automation is key to reproducibility.




2 Be mindful of the reviewer's time and effort.



3 The benefits of automation will also apply to you as the author!

Documentation





Artifact Documentation

```
1 # Overview
2 This repository contains the artifacts for the tutorial
3 "A Reproducible Tutorial on Reproducibility
4 in Database Systems Research".
5
6 ## Software Requirements
7 The artifact requires Docker and Docker Compose.
8
9 ## Hardware Requirements
10 At least 4GB of memory and 2 CPU cores are required.
11
12 ## Time Requirements
13 * Getting Started: 5 minutes
14 * Building the Docker image: ~20 minutes
15 * Running the experiments: ~5 minute
16
17 ## Reproducibility
18 To reproduce the experiments, follow these steps:
19 * Clone the repository
20 * In a terminal, navigate to the repository and run the following command:
21   ``bash
22   make run
23   ``
24 * The experiments will run and the results will be stored in the `data` directory.
25
26 ## Cleanup
27 To cleanup, run the following command:
28   ``bash
29   make clean
30   ``
```



03_advanced/README.md

- ▶ Clearly mention side effects of commands.
- ▶ Do not assume the user has any knowledge.
- ▶ Provide a step-by-step guide.
- ▶ Estimate the human time required for each step.
- ▶ Provide a cleanup script.
- ▶ **⚠** Do not break the user's system.

Concluding Remarks



Limitations of Containerization

- ▶ Not all software can be containerized.
- ▶ Not all workflows can be automated.
- ▶ Containerization provides the software environment, but not the hardware environment.
- ▶ Experiments may require specialized hardware.



Long-running Experiments, Specialized or Massive Hardware Requirements

- ▶ Few reviewers have a supercomputer at hand.
- ▶ Experiment should be executable on regular hardware (e.g., on a laptop).
- ▶ Provide reviewers with a downscaled version of the experiment.
 1. Use a smaller dataset,
 2. fewer iterations,
 3. fewer parameters.
- ▶ The main conclusions should be reproducible.

If not possible (which should be a rare exception), contact the program chairs for advice.



Cloud Deployment

- ▶ <https://docs.cloudlab.us/repeatable-research.html>
- ▶ <https://www.chameleoncloud.org/blog/2021/11/22/using-chameleon-for-artifact-evaluation/>

A Reproducible Tutorial on Reproducibility in Database Systems Research

VLDB 2024

Denis Hirn

{denis.hirn, tim.fischer}@uni-tuebingen.de

University of Tübingen

Tim Fischer

Gökhan Kul

gkul@umassd.edu

University of Massachusetts